

Spatial Econometrics

Dani Arribas-Bel

2016-04-17

This session¹ is based on the following references, which are good follow-up's on the topic:

- Session III of Arribas-Bel (2014). Check the “Related readings” section on the session page for more in-depth discussions.

This tutorial is part of Spatial Analysis Notes, a compilation hosted as a GitHub repository that you can access it in a few ways:

- As a download of a .zip file that contains all the materials.
- As an html website.
- As a pdf document
- As a GitHub repository.

Dependencies

The illustration below relies on the following libraries that you will need to have installed on your machine to be able to interactively follow along². Once installed, load them up with the following commands:

```
# Layout
library(tufte)
# For pretty table
library(knitr)
# Spatial Data management
library(rgdal)
# Pretty graphics
library(ggplot2)
# Pretty maps
library(ggmap)
# Various GIS utilities
library(GISTools)
# For all your interpolation needs
library(gstat)
# For data manipulation
library(plyr)
# Spatial regression
library(spdep)
```

¹ Points – Kernel Density Estimation and Spatial interpolation by Dani Arribas-Bel is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

² You can install package mypackage by running the command `install.packages("mypackage")` on the R prompt or through the Tools -> Install Packages... menu in RStudio.

Before we start any analysis, let us set the path to the directory where we are working. We can easily do that with `setwd()`. Please replace in the following line the path to the folder where you have placed this file -and where the `house_transactions` folder with the data lives.

```
setwd('/media/dani/baul/AAA/Documents/teaching/u-lvl/2016/envs453/code/GIT/kde_idw_r/')
#setwd('.')
```

Data

To explore ideas in spatial regression, we will be using house price data for the municipality of Liverpool. Our main dataset is provided by the Land Registry (as part of their Price Paid Data) but has been cleaned and re-packaged into a shapefile by Dani Arribas-Bel.

Let us load it up first of all:

```
hst <- readOGR(dsn = 'house_transactions', layer = 'liv_house_trans')

## OGR data source with driver: ESRI Shapefile
## Source: "house_transactions", layer: "liv_house_trans"
## with 6324 features
## It has 18 fields

## NOTE: rgdal::checkCRSArgs: no proj_defs.dat in PROJ.4 shared files
```

The tabular component of the spatial frame contains the following variables:

```
names(hst)

## [1] "pcds"      "id"        "price"
## [4] "trans_date" "type"      "new"
## [7] "duration"  "paon"      "saon"
## [10] "street"    "locality"  "town"
## [13] "district"  "county"    "ppd_cat"
## [16] "status"    "lsoa11"    "LSOA11CD"
```

The meaning for most of the variables can be found in the original Land Registry documentation. The dataset contains transactions that took place during 2014:

```
# Format dates
dts <- as.Date(hst@data$trans_date)
# Set up summary table
tab <- summary(dts)[c('Min.', 'Max.')]
tab
```

```
##           Min.           Max.
## "2014-01-02" "2014-12-30"
```

Although the original Land Registry data contain some characteristics of the house, all of them are categorical: *is the house newly built? What type of property is it?* To bring in a richer picture and illustrate how continuous variables can also be included in a spatial setting, we will augment the original transaction data with Deprivation indices from the CDRC at the Lower Layer Super Output Area (LSOA) level.

Let us read the csv in:

```
imd <- read.csv('house_transactions/E08000012.csv')
```

The table contains not only the overall IMD score and rank, but some of the component scores, as well as the LSOA code:

```
names(imd)

## [1] "LSOA11CD" "imd_rank" "imd_score"
## [4] "income"    "employment" "education"
## [7] "health"    "crime"      "housing"
## [10] "living_env" "idaci"      "idaopi"
```

That bit of information, LSOA11CD, is crucial to be able to connect it to each house transaction. To “join” both tables, we can use the base command `merge`, which will assign values from `imd` into `hst` making sure that each house transaction get the IMD data for the LSOA where it is located:

```
db <- merge(hst, imd)
```

The resulting table, `db`, contains variables from both original tables:

```
names(db)

## [1] "LSOA11CD" "pcds"      "id"
## [4] "price"     "trans_date" "type"
## [7] "new"       "duration"   "paon"
## [10] "saon"      "street"     "locality"
## [13] "town"      "district"   "county"
## [16] "ppd_cat"   "status"     "lsoa11"
## [19] "imd_rank"  "imd_score"  "income"
## [22] "employment" "education"  "health"
## [25] "crime"     "housing"    "living_env"
## [28] "idaci"     "idaopi"
```

Given there are 6,324 transactions in the dataset, a simple plot of the point coordinates implicitly draws the shape of the Liverpool municipality:

```
plot(db)
```



Figure 1: Spatial distribution of house transactions in Liverpool

Non-spatial regression, a refresh

Before we discuss how to explicitly include space into the linear regression framework, let us show how basic regression can be carried out in R, and how you can begin to interpret the results. By no means is this a formal and complete introduction to regression so, if that is what you are looking for, I suggest the first part of Gelman and Hill (2006), in particular chapters 3 and 4.

The core idea of linear regression is to explain the variation in a given (*dependent*) variable as a linear function of a series of other (*explanatory*) variables. For example, in our case, we may want to express/explain the price of a house as a function of whether it is new and the degree of deprivation of the area where it is located. At the individual level, we can express this as:

$$P_i = \alpha + \beta_1 NEW_i + \beta_2 IMD_i + \epsilon_i$$

where P_i is the price of house i , NEW_i is a binary variable that takes one if the house is newly built or zero otherwise and IMD_i is the IMD score of the LSOA where i is located. The parameters β_1 , β_2 , and β_3 give us information about in which way and to what extent each variable is related to the price, and α , the constant term, is the average house price when all the other variables are zero. The term ϵ_i is usually referred to as “error” and captures elements that influence the price of a house but are not whether the house is new or the IMD score of its area. We can also express this relation in matrix form, excluding subindices for i^3 .

Essentially, a regression can be seen as a multivariate extension of simple bivariate correlations. Indeed, one way to interpret the β_k coefficients in the equation above is as the degree of correlation between the explanatory variable k and the dependent variable, *keeping all the other explanatory variables constant*. When you calculate simple bivariate correlations, the coefficient of a variable is picking up the correlation between the variables, but it is also subsuming into it variation associated with other correlated variables –also called confounding factors⁴. Regression allows you to isolate the distinct effect that a single variable has on the dependent one, once we *control* for those other variables.

Practically speaking, running linear regressions in R is straightforward. For example, to fit the model specified in the equation above, we only need one line of code:

```
m1 <- lm('price ~ new + imd_score', db)
```

We use the command `lm`, for linear model, and specify the equation we want to fit using a string that relates the dependent variable

³ In this case, the equation would look like

$$P = \alpha + \beta_1 NEW + \beta_2 IMD + \epsilon$$

and would be interpreted in terms of vectors and matrices instead of scalar values.

⁴ **EXAMPLE** Assume that new houses tend to be built more often in areas with low deprivation. If that is the case, then *NEW* and *IMD* will be correlated with each other (as well as with the price of a house, as we are hypothesizing in this case). If we calculate a simple correlation between *P* and *IMD*, the coefficient will represent the degree of association between both variables, but it will also include some of the association between *IMD* and *NEW*. That is, part of the obtained correlation coefficient will be due not to the fact that higher prices tend to be found in areas with low IMD, but to the fact that new houses tend to be more expensive. This is because (in this example) new houses tend to be built in areas with low deprivation and simple bivariate correlation cannot account for that.

(price) with a set of explanatory ones (new and price) by using a tilde ~ that is akin the = symbol in the mathematical equation. Since we are using names of variables that are stored in a table, we need to pass the table object (db) as well.

In order to inspect the results of the model, the quickest way is to call summary:

```
summary(m1)

##
## Call:
## lm(formula = "price ~ new + imd_score", data = db)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -184254  -59948  -29032   11430 26434741
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)   235596      13326  17.679
## newY           4926       19104   0.258
## imd_score     -2416         308  -7.843
##              Pr(>|t|)
## (Intercept) < 2e-16 ***
## newY         0.797
## imd_score    5.12e-15 ***
## ---
## Signif. codes:
##  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 509000 on 6321 degrees of freedom
## Multiple R-squared:  0.009712, Adjusted R-squared:  0.009398
## F-statistic: 30.99 on 2 and 6321 DF, p-value: 4.027e-14
```

A full detailed explanation of the output is beyond the scope of this note, so we will focus on the relevant bits for our main purpose. This is concentrated on the Coefficients section, which gives us the estimates for the β_k coefficients in our model. Or, in other words, the coefficients are the raw equivalent of the correlation coefficient between each explanatory variable and the dependent one, once the polluting effect of confounding factors has been accounted for⁵. Results are as expected for the most part: houses tend to be significantly more expensive in areas with lower deprivation (an average of GBP2,416 for every additional score); and a newly built house is on average GBP4,926 more expensive, although this association can-

⁵ Keep in mind that regression is no magic. We are only discounting the effect of other confounding factors that we include in the model, not of *all* potentially confounding factors.

not be ruled out to be random (probably due to the small relative number of new houses).

Finally, before we jump into introducing space in our models, let us modify our equation slightly to make it more useful when it comes to interpreting it. Virtually every house price model in the literature is estimated in log-log terms:

$$\log P_i = \alpha + \beta_1 \log NEW_i + \beta_2 \log IMD_i + \epsilon_i$$

This allows to interpret the coefficients as *elasticities*, an economic term used to capture the percentual variation that a variable experiences as a result of a one percent increase in another one. This comes in very handy because it standardizes the results across variables. To fit such a model, we can specify the logarithm of a given variable directly in the formula. Note that we do not transform new, as it is a binary variable.

```
m2 <- lm('log(price) ~ new + log(imd_score)', db)
summary(m2)

##
## Call:
## lm(formula = "log(price) ~ new + log(imd_score)", data = db)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.3060 -0.3089 -0.0149  0.2936  5.3450
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)   13.54414    0.03504   386.56
## newY           0.23772    0.01934    12.29
## log(imd_score) -0.57471    0.01002   -57.36
##              Pr(>|t|)
## (Intercept)    <2e-16 ***
## newY           <2e-16 ***
## log(imd_score) <2e-16 ***
## ---
## Signif. codes:
##  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.516 on 6321 degrees of freedom
## Multiple R-squared:  0.3449, Adjusted R-squared:  0.3447
## F-statistic: 1664 on 2 and 6321 DF, p-value: < 2.2e-16
```

Looking at the results we can see a couple of differences with

respect to the original specification. First, the estimates are substantially different numbers. This is because, although they consider the same variable, they look at it from different angles, and provide different interpretations. For example, the coefficient for the IMD, instead of being interpretable in terms of GBP, the unit of the dependent variable, it represents an elasticity: a 1% increase in the degree of deprivation is associated with a 0.57% decrease in the price of a house.⁶ Second, the variable *new* is significant in this case. This is probably related to the fact that, by taking logs, we are also making the dependent variable look more normal (Gaussian) and that allows the linear model to provide a better fit and, hence, more accurate estimates. In this case, a house being newly built, as compared to an old house, is overall 23% more expensive.

⁶ **EXERCISE** How does the type of a house affect the price at which it is sold, given whether it is new and the level of deprivation of the area where it is located? To answer this, fit a model as we have done but including additionally the variable *type*. In order to interpret the codes, check the reference at the Land Registry documentation.

Spatial regression: a (very) first dip

Motivation

Spatial heterogeneity

- Spatial FE
- Spatial regimes

Spatial interaction

- Exogenous spatial effects
- Point to further spatial regression (lag, error) → Anselin (1988, 2003), plus Anselin & Rey (2015) for an in-depth treatment of modern approaches

Predicting house prices

- Show how to obtain a prediction for a given house
- Compare the estimate with that of interpolation?

References

Arribas-Bel, Dani. 2014. "Spatial Data, Analysis, and Regression-a Mini Course." *REGION 1* (1). European Regional Science Association: R1. http://darribas.org/sdar_mini.

Gelman, Andrew, and Jennifer Hill. 2006. *Data Analysis Using Regression and Multilevel/hierarchical Models*. Cambridge University Press.